

# Chapter 2

## **Working with Dates and Times**

## 2.1 Overview

In Python, there are core modules that contain base functionality to deal with dates and time, as well as the **pandas** and **NumPy** packages containing other methods. We will be focusing on the **pandas** methods for manipulating dates and times, using the **pandas** `Timestamp` object that is built upon the `NumPy datetime64` object.

## 2.2 Working with the `Timestamp` Object

An object of this type can be created using the `pd.to_datetime()` function on a string containing the date and time. The format of how the date time is constructed in the string can be done in a number of ways, where the function makes an appropriate guess as to what represents what in your date/time. The below strings all create the same `Timestamp` object.

```
>>> import pandas as pd
>>> pd.to_datetime("1999/01/12 22:01:00")
Timestamp('1999-01-12 22:01:00')
>>> pd.to_datetime("22:01 19990112")
Timestamp('1999-01-12 22:01:00')
>>> pd.to_datetime("199901122201")
Timestamp('1999-01-12 22:01:00')
>>> pd.to_datetime("12th Jan 1999 10:01pm")
Timestamp('1999-01-12 22:01:00')
```

Often our dates can be in more complex formats than those above. For example, the following should make the same time stamp as above, but doesn't:

```
>>> pd.to_datetime("12/01/1999 22:01")
Timestamp('1999-12-01 22:01:00')
```

We wish to have the 12th January date here, but have the 1st December returned. To amend this, we must use the `format` argument, which takes a string of `strftime` as follows:

```
>>> pd.to_datetime("12/01/1999 22:01", format="%d/%m/%Y %H:%M")
Timestamp('1999-01-12 22:01:00')
```



The format argument introduced above takes a strftime string, see <http://strftime.org/> for the code values of each directive.

### 2.2.1 Extracting Information from Timestamp Objects

Timestamp objects have an easy way of extracting characteristics in the form of attributes. These attributes are accessed using the dot "." and the attribute name; e.g. `Timestamp.hours`

```
>>> my_day = pd.to_datetime("12th Jan 1999 10:01pm")
>>> my_day.year

1999

>>> my_day.weekday_name

'Tuesday'
```

There are a large number of useful attributes for Timestamp objects, a full list of these attributes are detailed in the objects documentation <http://pandas.pydata.org/pandas-docs/stable/generated/pandas.Timestamp.html>.



When our data is stored in `Series` objects, we first need to use the `.dt` accessor to extract information.

## 2.3 The `pd.date_range` Function

Often, it is useful to create a sequence of dates, and this can be done via the `date_range` function from the `pandas` package. This creates an object of `DatetimeIndex` that contains a sequence of Timestamp objects.

```
>>> my_range = pd.date_range(start="1999-12-01", end="1999-12-09")
>>> my_range

DatetimeIndex(['1999-12-01', '1999-12-02', '1999-12-03', '1999-12-04',
               '1999-12-05', '1999-12-06', '1999-12-07', '1999-12-08',
               '1999-12-09'],
              dtype='datetime64[ns]', freq='D')
```

The default is to create a daily sequence but we can change that with either the `periods` argument or the `freq` argument or both.

```
>>> res = pd.date_range("19991201", "20000101", freq="4D")
>>> res

DatetimeIndex(['1999-12-01', '1999-12-05', '1999-12-09', '1999-12-13',
               '1999-12-17', '1999-12-21', '1999-12-25', '1999-12-29'],
              dtype='datetime64[ns]', freq='4D')

>>> pd.date_range(start="20000101", periods=4)

DatetimeIndex(['2000-01-01', '2000-01-02', '2000-01-03', '2000-01-04'],
              dtype='datetime64[ns]', freq='D')
```

## 2.4 Date Arithmetic

Dates and times are stored as numbers which means you can perform some arithmetic operations on them. For example, we can subtract two dates from each other.

```
>>> time_difference = pd.to_datetime("2018-01-31") -
pd.to_datetime("2018-01-01")
>>> time_difference

Timedelta('30 days 00:00:00')
```

This returns a `Timedelta` object which represents a period of time and is a subclass of `np.timedelta64` and `datetime.timedelta`. It shares some properties with the `Timestamp` which we can access in much the same way.

```
>>> time_difference.days

30

>>> time_difference.seconds

0
```

We can also create our own `Timedelta` object and add that to a `Timestamp` to do more complicated date arithmetic.

```
>>> newYearEve = pd.to_datetime("2017-12-31 23:59:59")
>>> newYearEve + pd.Timedelta("2 seconds")

Timestamp('2018-01-01 00:00:01')

>>> newYearEve + pd.Timedelta(3, unit='M')

Timestamp('2018-04-02 07:27:17')

>>> from pandas.tseries.offsets import *
>>> newYearEve - YearEnd(1)

Timestamp('2016-12-31 23:59:59')
```



See <http://pandas.pydata.org/pandas-docs/stable/timeseries.html#timeseries-offsets> for a full list of short hand versions of time spans.



1. Load in `dji.csv`, reformat the Date column, and store the weekday of that date in an additional column.
2. How many times does Monday occur in the data?
3. Load in `air_passengers.csv`. The Time column is in the format Month-Year. Reformat the column to a Day-Month-Year format.